

Readme for reproducibility submission of paper ID 338

A Source code info

Repository We provide one repository containing scripts for reproducing the experiments and two repositories containing source code. You only need to clone the first repository.

- The *reproducibility repository* is available under <https://github.com/jmuehlig/mxtasking-reproducibility>.
- The code repositories for (a) *MxTasking* and (b) comparative tree benchmarks using threads, Intel TBB, etc. are available under
 - (a) <https://github.com/jmuehlig/mxtasking>
 - (b) <https://github.com/jmuehlig/btree-benchmarks>.

Programming Language *C++* (17) for benchmarks, *Python* (3) for scripts, and *Java* for YCSB workload generation.

Compiler Info Clang in version *10.0.0-4ubuntu1*.

Packages/Libraries Needed You can verify all dependencies installed using the `check.sh` script.

If you are using *Ubuntu*, you can install all dependencies by executing `install_dependencies.sh`.

- CMake in version ≥ 3.10 for generating the Makefile
- Clang in version ≥ 10 for compilation
- Package `libnuma-dev`
- Python in version ≥ 3 for generating plots
- `curl` for downloading the YCSB workload
- Java for generating the YCSB workload
- Package `libtcmalloc-minimal4`
- Packages `libjemalloc-dev` and `libjemalloc2`
- Git for downloading the repositories
- Gnuplot for plotting
- Perf
- Intel VTune for cycle-based analysis

B Datasets info

YCSB We used the YCSB workload (in version 0.16.0) for the tree benchmarks. The script will download and build the workload automatically. The used workload specifications are `workloada` and `workloadc`, located in `workloads/ycsb`. The official repository is available at <https://github.com/brianfrankcooper/YCSB>. *This is just for your information, you do not have to change anything.*

TPC-H For the hash join experiment, we used the TPC-H benchmark with **scale factor** 100. The relevant files are `customer.tbl` and `orders.tbl`. **Notice:**

- Please download¹ and compile the TPC-H `dbgen` tool and set the directory containing the executable in `setup_environment.sh` (variable `DIR.TPCH.DBGEN`)
- or copy both `customer.tbl` and `orders.tbl` files into `workloads/tpch`.

If you set the `DIR.TPCH.DBGEN` variable, the script will automatically generate the needed files with scale factor 100 and move them into `workloads/tpch`.

C Hardware Info

Processor Two sockets with CPU *Intel(R) Xeon(R) Gold 6226 CPU @ 2.70GHz* (12 physical cores per socket, 24 logical cores per socket).

Caches The cache sizes for the entire system are specified.

- *L1d*: 768 KiB (32 KiB per physical core)
- *L1i*: 768 KiB (32 KiB per physical core)
- *L2*: 24 MiB
- *L3*: 38.5 MiB (19.25 MiB per socket)

Memory

- 6 DIMMs per socket (12 DIMMs total)
- Size: 15.26 GiB per DIMM (183 GiB total)
- Max Capable Speed: 2933 MHz

Secondary Storage Storage is not used.

Network Network is not used.

D Experimentation Info

We will describe how to use the scripts to reproduce the experiments, followed by a description of the files from the reproducibility project. In our experiments, we used Ubuntu 20.04 as the operating system.

¹http://tpc.org/tpc_documents_current_versions/download_programs/tools-download-request5.asp?bm_type=TPC-H&bm_vers=3.0.0&mode=CURRENT-ONLY

Instructions

Before running the experiments Please note before running the experiments.

- Please install all dependencies listed in A.
 - You might execute `check.sh` to verify all dependencies installed.
 - If you are running *Ubuntu* as the OS, you can execute `install_dependencies.sh` to automatically install missing dependencies, including Intel VTune.
- Please disable *NUMA balancing* by setting `/proc/sys/kernel/numa_balancing` to 0. *This will be done by the scripts, if you execute them as a root user.*
- Please set `/proc/sys/kernel/perf_event_paranoid` to -1 to allow the usage of (almost) all performance events. The performance counters are read by the executables. *This will be done by the scripts, if you execute them as a root user.*
- Please set `/proc/sys/kernel/kptr_restrict` to 0 when you are not executing the scripts as a root user.
- When executing the scripts as a root user, the scripts will enable CPU performance mode before running the benchmarks and disable it afterwards. *It is not required to run the scripts as root, however, for the experiments in our paper, we enabled the CPU performance mode.*

Running all at once You can run all experiments at once by executing `run_experiments.sh`. This will execute all the `run_*` files described next. However, please note the description for every single script below.

Running one experiment after the other You do not need to execute the following scripts, if you execute `run_experiments.sh`.

- The script `run_tree_experiments.sh` will produce Figure 10, Figure 11, and Figure 12 of the paper. Please note that the experiments might run for a long time (~ 36 hours in our setting).
- The script `run_hashjoin_experiments.sh` will produce Figure 9 of the paper. This should take about one hour.
 - **Please note:** You need to set `DIR_TPCH_DBGEN` (in `setup_environment.sh`) or generate the TPC-H files *customer.tbl* and *orders.tbl* on your own and copy them into `workloads/tpch`.
- The script `run_cycle_experiments.sh` will produce Figure 7 and Figure 13 of the paper.
 - **Please note:** You need to install *Intel VTune*² and set the variable `VTUNE_VARS_SCRIPT` in `setup_environment.sh` pointing to the *VTune* source script.

²<https://www.intel.com/content/www/us/en/develop/documentation/installation-guide-for-intel-oneapi-toolkits-linux/top/installation.html>

Results The plots will be stored in `plots/` and named like the Figures within the paper.

Configuration

There are some configurations that might be necessary depending on the underlying system. The default configuration should match the system described in C.

Library path Some executables need specific allocation libraries `tcmalloc` and `jemalloc`. You can specify their location by changing the value of `DIR_LIB` in `setup_environment.sh`.

Performance Counter We use different performance counter to better interpret the throughput-based results. One of them is the *stalled memory cycles* counter (`CYCLE_ACTIVITY.STALLS_MEM_ANY`). In our experience, this counter can be addressed differently on different architectures. The counter can be configured

- in `projects/mxtasking/src/benchmark/perf.cpp` (line 39) for `MxTasking`
- and in `project/tree-bench/src/util/perf.h` (line 104) for the tree benchmarks.

Description

The project includes scripts for downloading all sources and building executables needed for the experiments.

Shipped files Besides the mentioned scripts, the following files and folders are also included into the project.

- `workloads/` contains all files needed to build/run the workloads.
- `scripts/` contains some helper scripts for plotting and processing the raw output of the executables.
- `clear.sh` will remove all results, plots and downloaded sources.
- `check.sh` will check all dependencies and list missing libraries and packages.

Generated files The following files and folders are generated by running the scripts.

- `benchmark-results/` will contain the output generated by running the experiments and processed data (from the raw data) for plotting.
- `projects/` will contain the source code of (a) `MxTasking` and (b) a project for comparable tree experiments (thread-based `Blink-tree`, Intel TBB-based `Blink-tree`, `OLCTree`, open `BwTree`, and `Masstree`).
- `plots/` will contain all generated plots. The plots are named **Fig-*i*.pdf** where *i* is the number of the Figure within the paper.